

Proseminar Distributed Systems

Distributed Databases

(based on slides by Dr. Eva Zangerle)

January 12, 2021

- Estimates of International Data Corporation (IDC)
 - 2006: 0.18 zettabytes data (1 zettabyte = 1 billion terabytes)
 - 2011: 1.8 zettabytes
 - 2012: 2.7 zettabytes
 - 2018 → 2025: 175 zettabytes (49% in public cloud)
- Data from
 - Internet Archive, social networks (photos, posts)
 - LHC, NYSE, MRI, sensor networks, logs, personal footprint
- Problem
 - Access speeds have not kept up
 - New data needs, can't be squeezed into Relational Database Management Systems efficiently (e.g., graph data and traversal)

Big Data (cont.)

3

internetlivestats.com 12.01.2021 at around 8:00

- Twitter: 9,300 tweets
- Instagram: 1,043 photos
- Tumblr: 1,820 posts
- Skype: 5,,284 calls
- Internet traffic: 108,670 GB
- Google: 88,774 search queries
- Youtube: 87,682 videos
- Emails: 2,986,499 sent

internetlivestats.com 12.01.2021 at around 8:00

- Twitter: 9,300 tweets
- Instagram: 1,043 photos
- Tumblr: 1,820 posts
- Skype: 5,,284 calls
- Internet traffic: 108,670 GB
- Google: 88,774 search queries
- Youtube: 87,682 videos
- Emails: 2,986,499 sent
- Unit: per second

■ Aim of NoSQL

- Develop databases to target data amounts in terabyte / petabyte scale (web 2.0 age)
- Hard to scale relational systems with commodity hardware

■ Characteristics (mostly applicable)

- Non-relational, Schema-free
- Scale horizontally (scale-out)
- Open source (of course not strict: Amazon SimpleDB, Neo4j)
- Easy replication
- Consistency model BASE (often not ACID)
- Simple API
 - Complex queries often not possible
 - CRUD operators (create, read, update, delete)

- Avoid unneeded complexity
 - In some cases, no need for consistency, locking and logging
 - For Facebook, Twitter, etc. do we really need ACID?
 - Is "eventually consistency" enough?
- High throughput
- Horizontal scalability
- Cheap commodity hardware, Virtualization
- Trade off performance for reliability

- Key/value database
 - Database manager DBM (1979): stores data by use of a key in buckets (hashing)
- Document-oriented systems
 - IBM Lotus Notes (1984): stores user documents, groupware system
- Storage as Key/value pairs
 - Berkeley DB (1991)
- Column oriented systems
 - Sysbase IQ (1996)
- Distribution of these systems limited compared to RDBMS

- Databases from teh 70ies
 - Why popular now?
- Requirements changed
 - Semi-structure data in web age
 - Data volume increased
 - Amazon, google, Twitter, Facebook (MySQL, HBase)
 - Scalability
 - Independent data
 - Hard with distributed RDMS, expertise needed
 - ACID
 - Lots of different systems

Some Systems

8



chordless



■ **RE**remote **DI**ctionary **S**erver

- Key/value store
- In-memory (no data sets larger than memory)
- Can be persisted on disc
- Value include different data type

■ Use cases

- Caches in front of RDBMS
- github, flickr, twitter, stackoverflow

Sample Session (CLI)

10

```
SET england 44
EXISTS england // returns: 1
SET france 39
KEYS * // returns: england, france
DEL france
SET france 33
RANDOMKEY // returns: france
SET water cold
EXPIRE water 10 // expires in 10 seconds
GETSET water warm // returns: cold; set to warm
```

Expire

11

```
SET mykey "Hello"
```

```
OK
```

```
EXPIRE mykey 10
```

```
(integer) 1
```

```
TTL mykey
```

```
(integer) 10
```

```
SET mykey "Hello World"
```

```
OK
```

```
TTL mykey
```

```
(integer) -1
```

- Set up and a single Redis server on Amazon EC2 (via AWS SDK)
- Answer questions about fault tolerance in general and in Redis cluster
- Set up a Redis cluster on EC2 (via AWS SDK)
- Use Redis cluster to perform MapReduce